

CSCI-B 565 DATA MINING
Project Report
Fall 2016
Indiana University, Bloomington, IN

Ankit Saxena (ansaxena@iu.edu)
Samvat Rastogi (samrasto@iu.edu)
Shailendra Patil (patilsh@iu.edu)

December 15, 2016

Project 1: House Prices: Advanced Regression Techniques
Project 2: Outbrain Click Prediction

All the work herein is solely ours.

House Prices: Advanced Regression Techniques

Ankit Saxena, Samvat Rastogi, Shailendra Patil

Abstract

This project deals with predicting the Sales Price of individual residential properties in Ames, Iowa. The data set consists of 'Training Data' and 'Test Data'. The goal is to generate an appropriate regression model using the 'Training Data' and then predict the Sales Price for the given 'Test Data' based on the generated model. The 'Training Data' consists of 1460 observation and 81 Attributes including the Sales Price and 'Test Data' contains 1459 observations with 80 attributes which does not include Sales Price. Before applying any model, the data is preprocessed to take care of missing values and also to reduce number of attributes being used for model training. After data preprocessing, various regression methods like Linear Regression, Random Forest, Gradient Boosting, Extreme Gradient Boosting are applied and compared. Then, the model with the least error rate is used to predict the Sales Price of the houses using 'Test Data'.

Keywords

Data Mining — Kaggle — Housing — Regression — Prediction

Contents

Introduction	1
1 Data	1
1.1 Data Background	1
1.2 Data Description	1
1.3 Data Preprocessing	2
Missing Data Imputation • Outlier Analysis • Outlier Treatment • Feature Selection • Feature Engineering	
2 Algorithm and Methodology	4
2.1 Regression Techniques	4
2.2 Training of Regression Models	5
3 Experiments and Results	5
4 Summary and Conclusions	5
Acknowledgments	6
References	6

Introduction

"The Housing Market refers to the supply and demand for houses, usually in a particular country or region. A key element of the housing market is the average house prices and trend in house prices [1]. For many people buying a property is one of the important decisions in life. There are many factors like Area, Neighborhood, Connectivity, Transportation, House Price, etc. that impacts this decision.

Hence, it becomes very important to know the trend in which the houses are being sold, so that as a buyer we might have a rough idea what can be the expected price of a house. Predictive analysis is what is required here. As an analyst, we build various model and predict the rough estimate of the prices, which might be used by the property evaluators for their clients.

1. Data

1.1 Data Background

The data that has been provided for analysis consists of 3 files

1. **train.csv**
2. **test.csv**
3. **data_description.txt**

Training and Test data files consists of a column named ID, which can be used to uniquely identify every house and there are 79 features that can be used to predict the Sales Price of the house. Training data file consists of an additional column named Sale Price. The regression models have been developed on the training data to predict the value of Sales Price of the Test data. The data_description file provides all the necessary information regarding the attributes. The data was downloaded from Kaggle [2].

1.2 Data Description

The training data has 78 attributes excluding the ID and Sales Price and a total of 1460 records. The test data has 78 attributes excluding the ID and a total of 1459 records. Data consists of **43 categorical features** and **36 numerical features**.

The **categorical features** are: MSZoning, Street, Alley, LotShape, LandContour, Utilities, LotConfig, LandSlope, Neighborhood, Condition1, Condition2, BldgType, HouseStyle, RoofStyle, RoofMatl, Exterior1st, Exterior2nd, MasVnrType, ExterQual, ExterCond, Foundation, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, Heating, HeatingQC, CentralAir, Electrical, KitchenQual, Functional, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, PavedDrive, PoolQC, Fence, MiscFeature, SaleType, SaleCondition.

The **numerical features** are: MSSubClass, LotFrontage, Lot Area, OverallQual, OverallCond, YearBuilt, YearRemodAdd, MasVnrArea, BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, X1stFlrSF, X2ndFlrSF, LowQualFinSF, GrLivArea, BsmtFullBath, BsmtHalfBath, FullBath, HalfBath, BedroomAbvGr, KitchenAbvGr, TotRmsAbvGrd, Fireplaces, GarageYrBlt, GarageCars, GarageArea, WoodDeckSF, OpenPorchSF, EnclosedPorch, X3SsnPorch, ScreenPorch, PoolArea, MiscVal, MoSold, YrSold.

1.3 Data Preprocessing

The biggest task in any Data Mining problem is Data Preprocessing. Before we apply any algorithms, we first need to clean the data. The data might contain missing values, outliers, etc. There could be different kinds of attributes in the given data, like continuous variables, categorical variables, nominal, ordinal, ranges, ratios, etc. Such variables need to be taken care of based on the model we want to apply. Data preprocessing not only deals with these things but also covers feature selection, data reduction, etc. Given a set of attributes, not every attribute is important for data analysis or prediction. Few attributes might be of no use and few attributes will be very vital for our analysis. For this purpose we do feature selection.

1.3.1 Missing Data Imputation

There are many attributes which have NA (Not Applicable) as one of the accepted values in their domain. These NA's are not missing values but rather explain the absence of a particular feature in the house. The following attributes have NA in their domain: Alley, BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1, BsmtFinType2, FireplaceQu, GarageType, GarageFinish, GarageQual, GarageCond, PoolQC, Fence, MiscFeature.

The NA values for these attributes were replaced with NC while cleaning the data, to differentiate the actual missing values for other attributes with these 14 attributes.

After replacing the NA values with NC for these attributes, we found the actual number of NA values or missing values. A total of **357** values are missing in the Training data. A total of **358** values are missing in the Test data. The percentage of missing values in the training data is 0.3018%. The percentage of missing values in the test data is 0.3067%.

The majority of values are missing for the attribute LotFrontage. The percentage of missing values for LotFrontage in training data is 17.739%. The percentage of missing values for LotFrontage in test data is 15.558%. The mean of the values for attribute LotFrontage in training data is 70.05 and the median is 69.00. The mean of the values for attribute LotFrontage in test data is 68.58 and the median is 67.00. Thus, the missing values have been replaced with the mean of the existing values for LotFrontage, as there is not a huge difference between the mean and the median.

The attribute GarageYrBlt has the second highest number of missing values. The percentage of missing values for GarageYrBlt in training data is 5.547%. The percentage of missing values for GarageYrBlt in test data is 5.346%. The mean of the values for attribute GarageYrBlt in training data is 1979 and the median is 1980. The mean of the values for attribute GarageYrBlt in test data is 1978 and the median is 1979. Thus, the missing values have been replaced with the mean of the existing values for GarageYrBlt, as there is not a huge difference between the mean and the median.

Similarly for all other numerical attributes, the missing values were replaced by either mean or median based on the data and for categorical attributes the missing values were replaced by the mode value of that particular attribute.

1.3.2 Outlier Analysis

Outlier is an observation that diverges from the overall pattern in the sample. In any data, outliers affect the training model and sometimes this can be more than expected. Hence, such outliers should be taken care of.

The box plot for Sales Price is given below. As we can see,

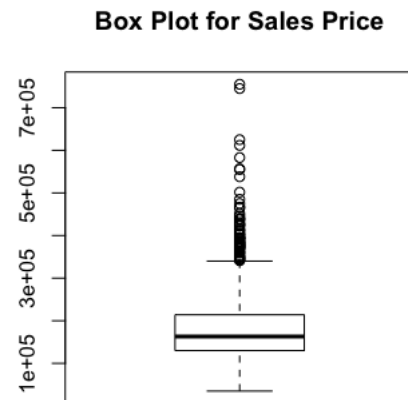


Figure 1.
Box Plot of Sale Price

there are a few outliers for the Sales Price, having a value greater than 500,000. We can compare the boxplot of Sales Price with some of the highly correlated categorical variables and check the outliers with respect to these categories.

1.3.3 Outlier Treatment

As outliers can affect the overall trend in the data, we will be removing these outliers from the training data. Here we have total 9 values of Sales Price which are greater than 500,000, which is just 0.61%. Hence, we can remove these rows from the data so that there would not be any value that greatly affects the predictions.




Figure 2.
Box Plot of Neighborhood




Figure 3.
Box Plot of Alley




Figure 4.
Box Plot of LandContour




Figure 5.
Box Plot of LotShape

1.3.4 Feature Selection

Feature Selection starts from an initial set of measured data and builds derived values (features) intended to be informative and non-redundant, facilitating the subsequent learning

and generalization steps, and in some cases leading to better human interpretations. Feature selection is related to dimensionality reduction. When the input data to an algorithm is too large to be processed and it is suspected to be redundant, then it can be transformed into a reduced set of features (also named a features vector). This process is called feature selection [3].

One of the methods for dimensionality reduction is Correlation. Correlation is a statistical technique that can show how strongly a pair of variables are related.

We are using Pearson's Coefficient to measure correlation and this can be done in the **R** tool using **cor** function. Below is the graph that shows the correlation among all the numerical attributes.




Figure 6.
Correlation Between Variables

As it is depicted in the graph, features like OverallQual, GrLiveArea, FullBath, GarageCars are among the few features which have high correlation with the Sales Price. But the Pearson's coefficient needs the data to be normal and linearly related. In our case, the size of the data is small and just by plotting the graphs it is not easy to predict if the data is normal and features have linear relationship or not. Hence, correlation is not a suitable technique for feature selection for our data.

The method that we have used for feature extraction is using **Boruta Package** in R tool. The algorithm is designed as a wrapper around a Random Forest classification algorithm. It iteratively removes the features which are proved by a statistical test to be less relevant than random probes. The **Boruta package** in **R** provides a convenient interface to the algorithm [4].

Below is the step wise working of **algorithm** [5]:

1. Firstly, it adds randomness to the given data set by creating shuffled copies of all the features (which are called shadow features).
2. Then, it trains a random forest classifier on the extended data set and applies a feature importance measure (the default is Mean Decrease Accuracy) to evaluate the importance of each feature where higher means more important.
3. At every iteration, it checks whether a real feature has a higher importance than the best of its shadow features (i.e. whether the feature has a higher Z score than the maximum Z score of its shadow features) and constantly removes features which are deemed highly unimportant.
4. Finally, the algorithm stops either when all features gets confirmed or rejected or it reaches a specified limit of random forest runs.

The Boruta package was used after the missing values were replaced, as Boruta does not work if we have missing values in data. After running the Boruta function and plotting the graph we get a list of **confirmed**, **tentative** and **rejected** features.




Figure 7.
Boruta Plot

The **green** box plots are for the features that have been confirmed, **yellow** for tentative and **red** for rejected features.

1.3.5 Feature Engineering

After cleaning the data and feature selection based on Boruta we have converted all the categorical attributes into dummy or indicator variables. These dummy variables are created based on the domain of each feature and store either 0 or 1 values, 1 if the value is present for that column or 0 if it is not present. For example, if a categorical feature Alpha has domain {A, B, C}, then Alpha would be replaced by three columns named A, B, C. For a particular row, if the value of Alpha is B, then A would be 0, B would be 1 and C would be 0.

These dummy variables tend to have values with more realistic interpretations.

2. Algorithm and Methodology

The project mainly deals with using various Regression techniques like **Linear Regression**, **Random Forest**, **Gradient**

Boosting, **Extreme Gradient Boosting** for predictive analysis.

What is Regression Analysis?

Regression analysis is a form of predictive modeling technique which investigates the relationship between a dependent (target) and independent variable(s) (predictor). This technique is used for forecasting, time series modeling and finding the causal effect relationship between the variables [6].

2.1 Regression Techniques

1. Linear Regression

A linear regression establishes a relation between a dependent variable (Y) and one or more independent variable(X) using a best fit straight line (also known as "Regression Line").

It is represented by below equation

$$y = a + bx \quad (1a)$$




Figure 8.

Linear Regression Line

2. Random Forest Regression

Random forests or random decision forests are an ensemble learning method for classification, regression and other tasks, that operate by constructing a multitude of decision trees at training time and outputting the class that is mean prediction (regression) of the individual trees [7].

3. Gradient Boosting

Gradient boosting is a machine learning technique for regression and classification problems, which produces a prediction model in the form of an ensemble of weak prediction models, typically decision trees. It builds the model in a stage-wise fashion like other boosting methods do, and it generalizes them by allowing optimization of an arbitrary differentiable loss function [8].

4. **Lasso Regression:** Lasso (least absolute shrinkage and selection operator) is a regression analysis method that performs both variable selection and regularization in order to enhance the prediction accuracy and interpretability of the statistical model it produces.

2.2 Training of Regression Models

The method employed to train a prediction model is as follows:

1. The given training data is partitioned into two parts. 80% of the given training data was considered as training data for initial model training and remaining 20% was considered as test data, for the purpose of calculating the error rates.
2. We use this 80% data on various regression models.
3. The model was applied on the 20% of the data used as test data, and the Sales Price was predicted.
4. The RMSE was calculated using the actual Sales Price and predicted Sales Price.
5. This process was repeated on all regression models.

A 10-fold cross-validation was also applied to decide the optimal number of trees. The figure shows the number of trees used and the average error. We have found that as the number of trees increases, the average error decreases until it converges and then starts increasing again. Other studies have also concluded the same results, that there is a threshold value beyond which there is no advantage in increasing the number of trees.




Figure 9.

Random Forest 10-fold cross-validation

Here are some of the parameters that we have optimized for our Random Forest model on the training data:

Number of predictors sampled for splitting at each node: 65
Number of trees: 500

Similarly, for the Gradient Boosting model we have found the optimal values for some parameters on the training data:

Number of trees to fit: 10000
Maximum depth of variable interaction: 15
Learning Rate or step-size reduction: 0.002
Minimum number of observations in tree terminal nodes: 15

3. Experiments and Results

All the experiments were performed on a machine with two Intel Xeon E5-2650 v2 8-core processors with 32GB of mem-

ory, available at Indiana University, Bloomington. We have used R version 3.3.0, R Studio version 0.99.902 for the coding involved in this project.

We have compared the regression models used for predicting the Sales Price of the houses based on the root mean square error (RMSE). Root Mean Square Error is a measure of the quality of an estimator and is the difference between the predicted and observed values.

The formula used for calculating RMSE = $\sqrt{\frac{(\text{predicted} - \text{observed})^2}{n}}$

After training the regression models, the following values were calculated for RMSE:

1. Linear Regression: 33262.7
2. Simple Trees: 47717.98
3. Random Forest: 28322.66
4. Gradient Boosting: 25877.12
5. Lasso: 27652.61
6. Extreme Gradient Boosting: 30381.65

Also, **Ensemble Learning** was used to predict the Sales Price of the houses. Ensemble learning creates a predictor or model by combining multiple predictors (from Linear Regression, Random Forest, etc.) to create a stronger overall prediction. An ensemble with two techniques that are more diverse might give much better results than the one with similar techniques involved. The prediction models for Gradient Boosting and Lasso regression were combined to create the ensemble model here. The RMSE value for this Ensemble model was calculated to be 21887.75.

After calculating the RMSE values we found out that the ensemble model has the least RMSE among all the predictors used.

The experiment was repeated on the entire training data to train the models again and the Ensemble model was used to predict the final House Prices, as it has the least RMSE out of all the prediction models.

4. Summary and Conclusions

The main challenge faced in predicting the Sales Price of the houses was to train the prediction models and find the optimal parameters for each model.

The large number of observations and the attributes present in the training dataset have helped us to analyze and predict the Sales Price of houses correctly, to an extent. We have gained insight about various phases of the Data Mining processes like data preparation, transformation and its validation. It has been found out that a lot of the variation in the Sales Price can be explained by observing the features like GrLivArea, GarageArea, OverallQual, and TotalBsmtSF.

For future work, we would like to train the Extreme Gradient Boosting model more efficiently and also implement other techniques like Ridge Regression, Neural Network.

Acknowledgments

We have put in considerable efforts to reach this stage of the project but it would not have been possible without the continuous guidance and support of many other individuals. We would like to extend our sincere thanks to all of them.

We are highly indebted to our Professor Mehmet Dalkilic, for his supervision as well as for providing the necessary resources required for the completion of this project report.

We would also like to thank our Associate Instructor Hasan Kurban for his help and constant support throughout our coursework.

References

- [1] Tejvan Pettinger. *Definition of the housing market*. Economics Help, 2015.
- [2] *Data - House Prices: Advanced Regression Techniques*. Kaggle, 2016.
- [3] *Feature extraction*. Wikipedia.
- [4] Witold R. Rudnicki Miron B. Kursa. *Feature Selection with the Boruta Package*. Journal of Statistical Software, 2010.
- [5] Debarati Dutta. *How to perform feature selection (i.e. pick important variables) using Boruta Package in R?* Analytics Vidhya, 2016.
- [6] Sunil Ray. *7 Types of Regression Techniques you should know!* Analytics Vidhya, 2015.
- [7] *Random forest*. Wikipedia.
- [8] *Gradient boosting*. Wikipedia.

Outbrain Click Prediction

Ankit Saxena, Samvat Rastogi, Shailendra Patil

Abstract

In today's world, Internet is omni-present. From finding a meaning of a simple word to operating a remote machine, Internet helps us in all. As we are advancing in the technology, our need from the technology is changing day by day. We are using technology to make our daily task simpler to more intelligent. One of those task is content recommendation. A content recommendation system is an intelligent system which helps to recommend appropriate content to the users based on their activities like browsing patterns on Internet. Outbrain is one such organization which pairs the relevant content with readers in about 250 billion personalized recommendation every month.

The purpose of our project is to help Outbrain to improve its recommendation system by predicting which contents are likely to be clicked by its global user base. Hence, as our goal, we created a prediction model using the train data to predict whether the user will click a particular content. We trained the model with train data with important attributes like platform, advertisement id and display id from meta data to complete the prediction. Since data is relational in nature, these huge data sets could be joined on a basis of single key resulting into data set with required data or with required features. For the model, we calculated the probability of an advertisement to get clicked from a particular platform on the basis of which we could determine the likeliness of the click on the advertisement.

Keywords

Data Mining — Kaggle — Prediction — Probability

Contents

Introduction	1
1 Background	1
1.1 Data	1
Data Fields • Data Preprocessing	
2 Algorithm and Methodology	2
2.1 Training of the model	2
3 Experiment and Result	3
3.1 Hardware used	3
3.2 Experiment	3
3.3 Results	3
4 Summary and Conclusions	3
4.1 Future Work	3
Acknowledgments	3
References	4

Introduction

Today, marketing and selling a product is far more easier than old days. Thanks to Internet. Due to Internet, a product can be advertised and purchased from any corner of the world in the matter of few seconds. Moreover, the product in the advertisement can be anything from a simple pen to an insurance policy. This means anything can be advertised on the Internet, be it a product or a service. But the major

challenge is to advertise the correct product or service to the target user. Because advertising a shampoo to a bald guy will never help the shampoo company to make a profit. So, how can this accuracy be achieved?

Outbrain is an advertising company which uses behavioral targeting to recommend articles, products, services, photos, videos to the user. Behavioral Targeting is a technique which allows Outbrain to increase the efficiency of its advertisement network by user web browsing behavior information. According to Wikipedia [1], "behavioral targeting uses information collected from an individual's web-browsing behavior (e.g., the pages that they have visited or searched) to select advertisements to display". Because of this technique, Outbrain's recommendations can be found on more than 35,000 websites and it serves over 250 million recommendations and 15 billion page views per month. It reaches over 87% of US Internet users.

1. Background

The problem has been picked up from Kaggle. The challenge stated in the problem is to predict which advertisement are likely to be clicked.

1.1 Data

As this is a Kaggle Problem [2], the dataset was downloaded from Kaggle [3]. The dataset contains a sample of users' page views and clicks on multiple publisher sites in the United

States between June 14th June, 2016 to 28th June, 2016. It contains numerous sets of content recommendations served to a specific user in a specific context. Each context (i.e. a set of recommendations) is given a `display_id`. In each such set, the user has clicked on at least one recommendation. The identities of the clicked recommendations in the test set are not revealed. This is a very large relational dataset.

1.1.1 Data Fields

Each user is represented by a unique id, namely `uuid`. A user can view any web page with content, which is denoted by an id (`document_id`). Each document has a set of ads (`ad_id`). Each of these ads belong to a campaign (`campaign_id`) run by an advertiser (`advertiser_id`). Dataset also provides meta data about the document like mentioned entities, categories, topics etc. Files used for doing further computation are:

1. `events.csv`
2. `clicks_train.csv`
3. `clicks_test.csv`

`clicks_train.csv` is the training set, showing which set of ads was clicked. This file has total of 87141731 rows. It has no missing values.

Table 1. `clicks_train.csv`

Column	Distinct Values	Missing	Data Type
<code>display_id</code>	16874593	0	text
<code>ad_id</code>	478950	0	text
<code>clicked</code>	2	0	text

`clicks_test.csv` is the same as `clicks_train.csv`, except it does not have the `clicked` ad. Each `display_id` has only one `clicked` ad. This dataset contains `display_ids` from the entire dataset time frame. It has a total of 32225162 rows with 6245533 distinct `display_ids` and 381385 distinct `ad_id`.

Table 2. `clicks_test.csv`

Column	Distinct Values	Missing	Data Type
<code>display_id</code>	6245533	0	text
<code>ad_id</code>	381385	0	text

`events.csv` provides information on the `display_id` context. It covers both the train and test set. This dataset has 23120126 rows.

Table 3. `Events.csv`

Column	Distinct Values	Missing	Data Type
<code>display_id</code>	23120126	0	text
<code>uuid</code>	19794967	0	text
<code>document_id</code>	894060	0	text
<code>timestamp</code>	22896622	0	text
<code>platform</code>	4	0	text
<code>geo_location</code>	2988	801	text

1.1.2 Data Preprocessing

Data Preprocessing plays a very important role in any data mining problem. This problem have real world data which is collected from users from various geographies. Since data is huge in size, it is highly possible that some of the data is missing or distorted. Table 1-3 shows an overview of Data. There are no missing values for `click_train.csv` and `clicks_test.csv`. However, there are 340 missing values and 461 blank values in `events.csv` for `geo_location`. But since the data is very huge and total defective rows are much less than 0.05% of total data, we can remove those rows to sanitize the data.

It is quite noticeable that each data set has very limited number of features and are dependent on the features available in different data set. Moreover, data is relational in nature. In the sample submission file, only `display_id` and `ad_id` are required. So, in order to obtain the appropriate output, `clicks_train` and `events` data were merged, on the basis of `display_id` and `platform`. This means that data is joined using **Left Outer join**. The feature, `platform`, was used because browsing device can actually change the browsing behavior of a user as some websites are not mobile compatible or optimized.

2. Algorithm and Methodology

This problem mainly deals with prediction of user clicks. Hence, to predict the same, we implemented the solution using **regularized probabilities**.

Regularization refers to a process of introducing additional information in order to solve an ill-posed problem or to prevent over fitting. Mathematically, a regularization term $R(f)$ is introduced to a general loss function:

$$\min_f \sum_{i=1}^n V(f(\hat{x}_i), \hat{y}_i) + \lambda R(f)$$

for a loss function V that describes the cost of predicting $f(x)$ when the label is y , such as the square loss or hinge loss, and for the term λ which controls the importance of the regularization term. $R(f)$ is typically a penalty on the complexity of f , such as restrictions for smoothness or bounds on the vector space norm [4].

For this, we combined the `ad_id` from `clicks_train` and `platform` from `events` into a single key, and the probability for each `ad_id` corresponding to a `display_id` was calculated. And we also calculate the probability of total advertisement clicked, so that if at all there is any `ad_id` which was not trained in the training data, then we make the probabilities of such entries equal to the probability of advertisements clicked. This is known as Smoothing. Using these probabilities we predict the probabilities for the given `display` and `ad_id` in test data.

2.1 Training of the model

In order to train the model, we selected the `events` and `clicks_train` data sets and merged them using a left outer join. Then, the columns of `ad_id` and `platforms` are joined together

in order to form a unique key, say `ad_id_pl`. Now, the probability of all `ad_id` corresponding to its `display_id` is calculated and also the total probability of ads being clicked i.e. 0 or 1 is calculated. For this training data, the total probability of ad being clicked is 0.1936454 as there are 16874593 advertisements that were clicked. These probabilities were calculated using regularization, where we tuned the model by applying penalty for extreme cases.

Further, these calculated probabilities were used to predict the probabilities of the test data corresponding to their `ad_id` and if there is any untrained data point present then 0.1936454 was used as its probability. Hence, as per the requirement of Kaggle, the `ad_id` obtained were sorted in ascending order according to its probability and then grouped by `display_id`.

3. Experiment and Result

3.1 Hardware used

All the experiments were performed on a machine with two Intel Xeon E5-2650 v2 8-core processors with 32GB of memory, available at Indiana University, Bloomington. We have used R version 3.3.0, R Studio version 0.99.902 for the coding involved in this project.

3.2 Experiment

As data was provided by Kaggle, the complete data set is divided into two major parts of train and test. The train data consists of all the `ad_id` with respective `display_id` and a click flag. The test data consists of same type of data apart from the clicked flag.

The standard technique used by Kaggle, Mean Average Precision (MAP) for precision evaluation is being used to compare the models.

Mean Average Precision (MAP) is a number metric which gives the measure of performance of an algorithm. MAP is an average of precision values of all points where relevant information can be found. The average precision values at the ranks of relevant information are found. This measure is very different from precision as precision only considers a single threshold in the rank and average precision considers the entire ranking, assigning it a quality number to it. But while executing, it gives more importance to the things which happens in the early stages of ranking. This means that the information retrieved at rank 1 is much more important than rank 2 [5].

Mathematically, average precision is defined as below:

$$ap@n = \sum_{k=1}^n \frac{P(k)}{\min(m, n)}$$

where m are the total edges in any graph and n is the maximum number of nodes, a user is likely to guess. And, The mean average precision for N users at position n is the average of the average precision of each user, i.e.,

$$MAP@n = \sum_{i=1}^N ap@n_i / N$$

Here, we have leveraged *MAP@12* algorithm in Kaggle, which is defined as

$$MAP@12 = \frac{1}{|U|} \sum_{u=1}^{|U|} \sum_{k=1}^{\min(12, n)} P(k)$$

where $|U|$ is the number of `display_ids`, $P(k)$ is the precision at cutoff k , n is the number of predicted `ad_ids` [6] [7].

3.3 Results

On evaluating, the performance of our algorithm with regularization parameter of 18, we received a MAP score of **0.63853** and initial ranking of **121**.

4. Summary and Conclusions

The given data was tremendously huge and required proper prior knowledge before taking any step. Hence it is very important to study the data very well. We were able to predict the likelihood of user clicks using very less features to a certain extent of accuracy. This project helped us a lot in gaining insight about handling large set of data, importance of infrastructure dependability and preprocessing of data. In the end, we were able to secure a score of **0.63853** and initial ranking of **121** (As on Dec 14th, the rank was 189).

4.1 Future Work

We would like to extend this project further till the actual deadline on Kaggle. As we mentioned, only few features were used to create model. Hence, there is possibility to create more accurate model using other features. We would further like to consider creating model based on logistic regression, SVM and neural network as we believe we can achieve better results with those techniques.

Acknowledgments

We have put in considerable efforts to reach this stage of the project but it would not have been possible without the continuous guidance and support of many other individuals. We would like to extend our sincere thanks to all of them.

We are highly indebted to our Professor Mehmet Dalkilic, for his supervision as well as for providing the necessary resources required for the completion of this project report.

We would also like to thank our Associate Instructor Hasan Kurban for his help and constant support throughout our coursework.

References

- [1] *Behavioral targeting*. Wikipedia.
- [2] *Outbrain Click Prediction*. Kaggle, 2016.
- [3] *Data - Outbrain Click Prediction*. Kaggle, 2016.
- [4] *Regularization (mathematics)*. Wikipedia.
- [5] Victor Lavrenko. *Mean Average Precision*. Youtube, 2014.
- [6] Mikel Bober. *Mean Average Precision*. Kaggle, 2016.
- [7] *Outbrain Click Prediction: Evaluation*. Kaggle, 2016.